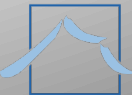




Informatics II for Engineering Sciences (MSE)

Chapter IV – Databases: SQL



Normal forms - Motivation

- **Normalization** of schema by **splitting attributes** into several **relations**
- Splitting has to be without loss of **information**
- Goal:
 - Remove unnecessary redundancy
 - ➔ Avoid inconsistency when data is changed
 - ➔ Reduce memory consumption
 - Minimize redesign when extending the database structure
 - Allow general queries



Normal forms

- Normal forms sorted by strictness
 - **1st Normal form:** Each attribute has atomary domain

ProfessorLecture					
<u>PersNo</u>	Name	Rank	Room	<u>Title</u>	SWS
2125	Sokrates	C4	226	{5041, Ethics}	4
2125	Sokrates	C4	226	{5049, Maieutics}	2
2125	Sokrates	C4	226	{4052, Logic}	4
...

Normal forms

- Normal forms sorted by strictness
 - **1st Normal form:** Each attribute has atomary domain
 - **2nd Normal form:** No attribute that is not a key is functionally dependent of a subset of any key candidate

ProfessorLecture					
<u>PersNo</u>	Name	Rank	Room	<u>Title</u>	SWS
2125	Sokrates	C4	226	Ethics	4
2125	Sokrates	C4	226	Maieutics	2
2125	Sokrates	C4	226	Logic	4
...

Normal forms

- Normal forms sorted by strictness
 - **1st Normal form:** Each attribute has atomary domain
 - **2nd Normal form:** No attribute that is not a key is functionally dependent of a subset of any key candidate
 - **3rd Normal form:** No attribute that is not a key has a transitive dependency from a key candidate

CD			
CDID	Albumtitle	Singer	Founding year
4711	Not that kind	Anastacia	1999
4712	Wish you were here	Pink Floyd	1964
4713	Freak of Nature	Anastacia	1999
...



Normal forms

- Normal forms sorted by strictness
 - **1st Normal form:** Each attribute has atomary domain
 - **2nd Normal form:** No attribute that is not a key is functionally dependent of a subset of any key candidate
 - **3rd Normal form:** No attribute that is not a key has a transitive dependency from a key candidate
 - Boyce-Codd Normal form
 - 4th Normal form
 - 5th Normal form
- A Normal form fullfills all criteria of the previous normal forms



SQL

- **Structured Query Language**
- Standardized language for
 - ➔ Making Queries
 - Data definition (DDL)
 - Data manipulation (DML)
- The language is standardized, however most implementations do not follow the standard entirely

You can try out the queries shown in the lecture on:

<http://www.hyper-db.com/interface.html>

SQL-Query

select personalNr, name ——— **attributes**

from professor ——— **relations**

where rank = 'C4';

condition

Professor	
<u>persNo</u>	name
2125	Sokrates
2126	Russel
2136	Curie
2137	Kant

- Boolean operators: and, or, not
- Sort result: order by, asc, desc
- Eliminate duplicates: distinct
- Set operators: union, intersect, minus
- Set comparisons: exists, all, in
- Aggregate functions: avg, max, min, count, sum
- Subqueries

SQL-Queries: Syntactical Sugar

```
select *  
from Student  
where semester > = 1 and semester < = 4;
```

```
select *  
from Student  
where semester between 1 and 4;
```

```
select *  
from Student  
where semester in (1,2,3,4);
```

SQL-Queries: Syntactical Sugar

```
select *  
from Student  
where name like `T%eophrastos`;
```

```
select distinct s.name  
from Course c, participates p, Student s  
where s.matrNo = p.matrNo and  
      p.courseNo = c.courseNo and  
      c.title like `%thik%`;
```

SQL-Queries: Syntactical Sugar

Place holder: "%" ; "_"

- "%" for any symbols or no symbol
- "_" for exactly one symbol

select *

from Student

where name **like** `T%eophrastos`;

select distinct s.name

from Course c, participates p, Student s


where s.matrNo = p.matrNo **and**

p.courseNo = c.courseNo **and**


c.title **like** `%thik%`;

SQL-Queries: Joins


SQL JOINS




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```




```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```

© C.L. Moffatt, 2008

SQL

- **Structured Query Language**
- Standardized language for
 - Making Queries
 -  - Data definition (DDL)
 - Data manipulation (DML)
- The language is standardized, however most implementations do not follow the standard entirely

You can try out the queries shown in the lecture on:

<http://www.hyper-db.com/interface.html>

SQL-Queries: Schema definition - Create database

create table Student

(matrNo **integer primary key**,

name **varchar(30) not null**,

semester **integer check** semester)

Student		
<u>matrNo</u>	name	semester

SQL-Queries: Schema definition - Create database with Composite Primary Key

```
Create table participates (  
    matrNo int(11) not null,  
    courseNo int(11) not null,  
    primary key (matrNo, courseNo)  
)
```

participates	
<u>matrNo</u>	<u>courseNo</u>

SQL-Queries: Data integrity

Foreign Key

- Point towards a tuple in another relation
e.g. taughtBy in Course relates to the primary key in the relation Professor
- Referential integrity: Foreign keys need to refer to another existing tuple or contain a null value

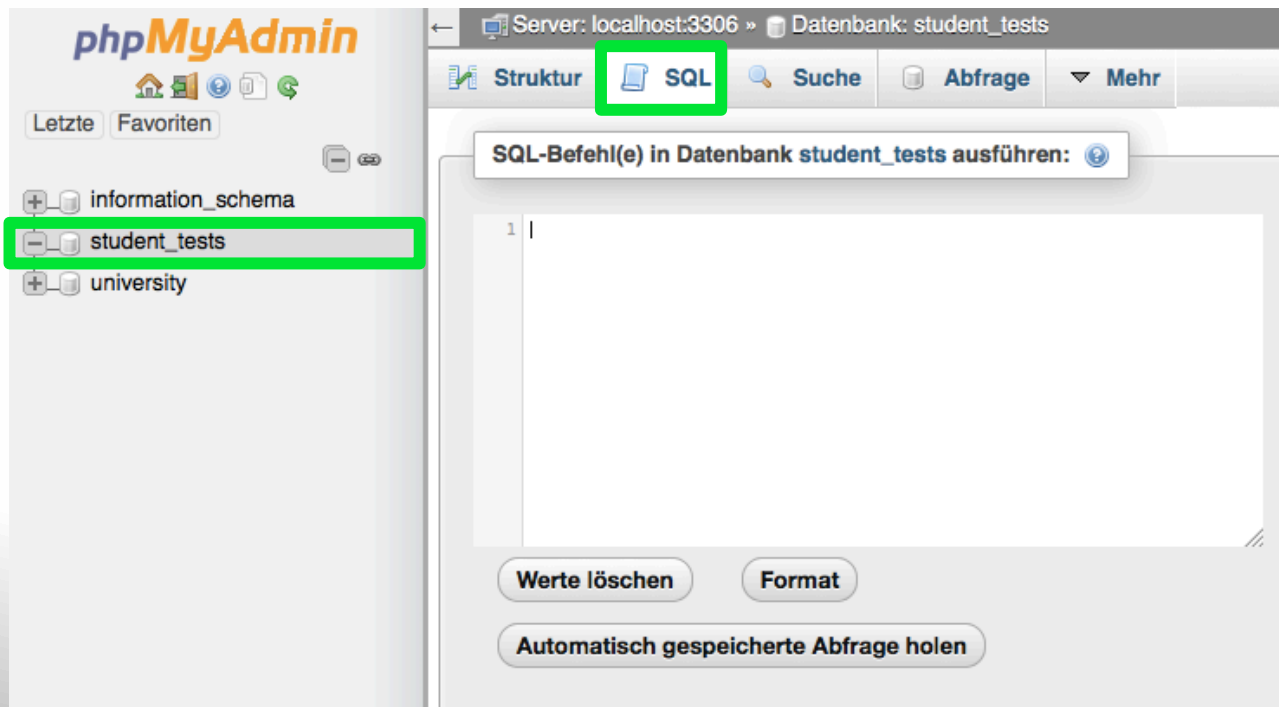
Course			
<u>courseNo</u>	title	sws	taughtBy

create table Vorlesungen

(courseNo	integer primary key,
title	varchar(30),
sws	integer,
taughtBy	integer references Professoren on delete set null);

Exercise 1

- Open <https://infeedbruegge.in.tum.de/mse> in browser
 - Username: mse
 - Password: mse
- Open the database **student_tests** and switch to **SQL** tab



Exercise 1 – Create a table...

- Named “*YourNameContacts*” where you replace *YourName* with your actual name
- With the following attributes:
 - contactID, integer, not null
 - firstName, varchar, length 25, not null
 - lastName, varchar, length 25, not null
 - phoneNumber, varchar, length 25, not null
- With composite primary key, consisting of
 - contactID
 - phoneNumber




Exercise 1 - Solution

```
create table BarbaraReichartContacts (  
    contactID integer not null,  
    firstname varchar(25) not null,  
    lastname varchar(25) not null,  
    phoneNumber varchar(25) not null,  
    primary key (contactID, phoneNumber)  
);
```



SQL

- **Structured Query Language**
- Standardized language for
 - Making Queries
 - Data definition (DDL)
 -  - Data manipulation (DML)
- The language is standardized, however most implementations do not follow the standard entirely

You can try out the queries shown in the lecture on:

<http://www.hyper-db.com/interface.html>

SQL – Changing database state

Insert tuples

insert into participates

select matrNo, courseNo

from Student, Course

where title = `Logic` ;

insert into Studenten (matrNo, name)

values (28121, `Archimedes`);

Student		
matrNo	name	semester
⋮	⋮	⋮
29120	Theophrastos	2
29555	Feuerbach	2
28121	Archimedes	-

SQL – Changing database state

Delete Tuples

delete Student

where semester > 13;

Change Tuples

update Student

set semester = semester + 1;

Exercise 2 – Insert Contacts

- Open the table “*YourNameContacts*”
- Add the following entries:

YourNameContacts			
<u>contactID</u>	firstName	lastName	<u>phoneNumber</u>
1	Jan	Knobloch	0176/173 143 21
1	Jan	Knobloch	0175/223 333 44
2	Emitza	Guzman	089/289 18213
4	Bernd	Brügge	
3	Barbara	Reichart	089/289 18213
3	Alfons	Reichart	089/289 18213

- You cannot create some of the entries due to constraints, which ones? What is the constraint that hinders creation?

Homework #9 - SQL


- Write the following queries for the **university** database:
 - Select minimal **and** maximal semester of all students
 - Select all students that take a class taught by Sokrates. Be sure to eliminate possible duplicates (e.g. when Susan visits two of Sokrates classes)
- Edit tables in **student_tests**
 - Create a new table that contains phone number and type (e. g. mobile, work, private, ...). You can store the type as a varchar
 - Make sure phone number is a foreign key
 - Add 5 tuples, then remove one of them
- Test your queries in <https://infeedbruegge.in.tum.de/mse>
 - Username: mse
 - Password: mse






Informatics II for Engineering Sciences (MSE)

Chapter III – Using Databases from Java



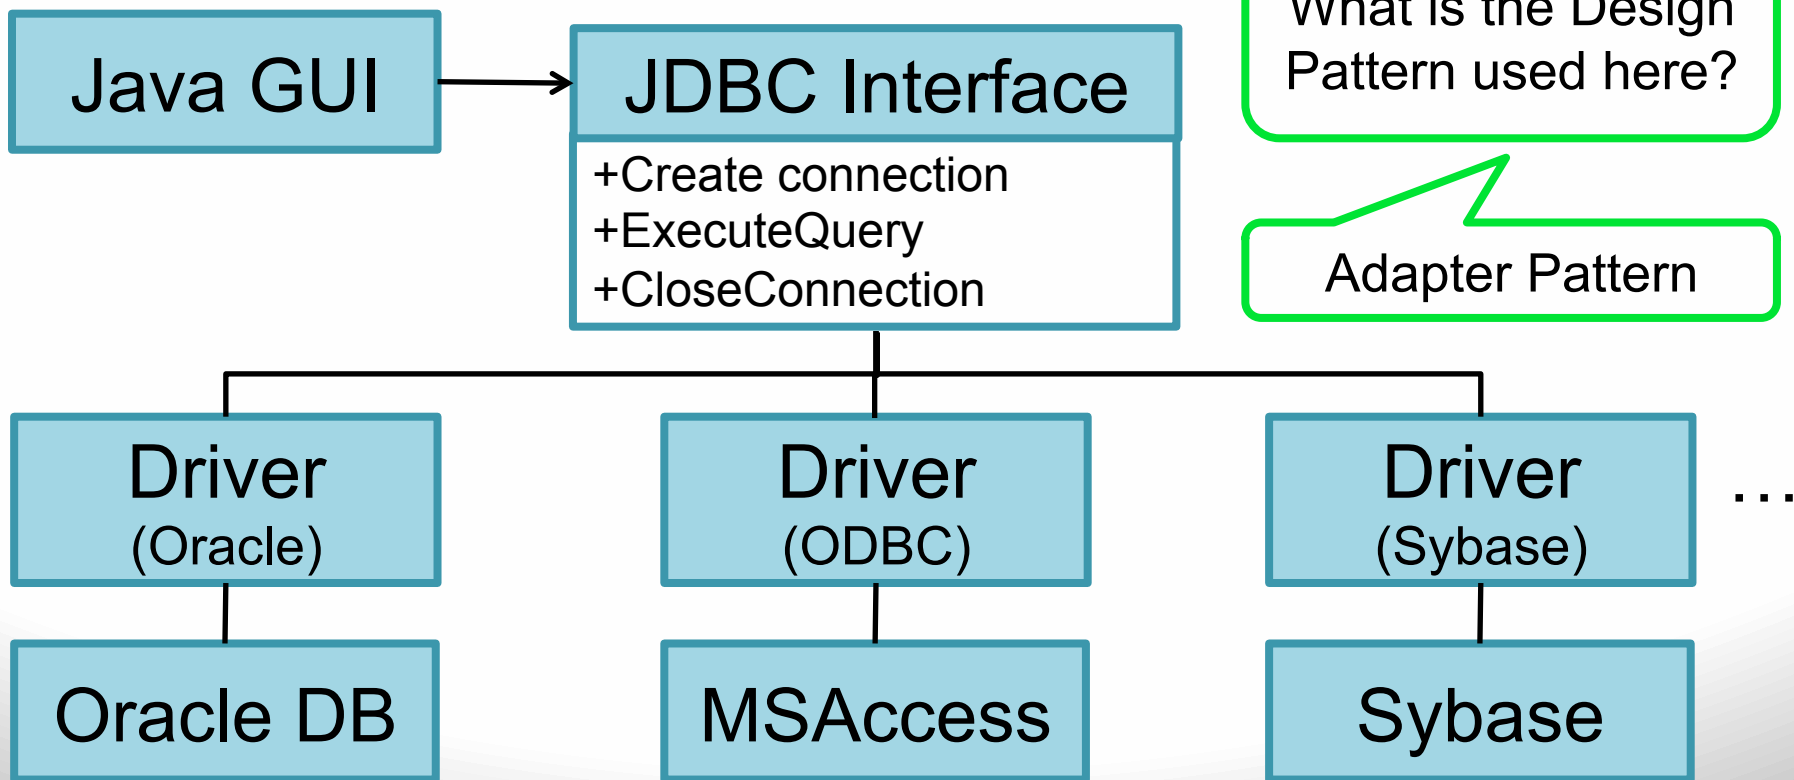
Use Database in Java Code

- We showed you how to map UML to Java
- We showed you how to map UML to Tables
- How do we connect Java Code with Tables?



Interfacing Java to Relational Database: JDBC

- Standardized interface to connect relational databases with Java



JDBC: Create Connection

```
Statement statement = null;
Connection connection = null;
try {
    Class.forName("org.h2.Driver");
    connection = DriverManager.getConnection("jdbc:h2:~/
        info2", "sa", "");
    statement = connection.createStatement();
} catch (Exception e) {
    System.err.println("Error: " + e);
    System.exit(-1);
}
```

Load
Driver

Create
statement

Exception
handling

Open
connection
using:

- Database url
- User name
- Password

JDBC: Execute query

```
try {  
    ResultSet resultSet = statement.executeQuery("select  
        avg(semester) from Student");  
    resultSet.next();  
    System.out.println("Average age: " +  
        resultSet.getDouble(1));  
    resultSet.close();  
} catch (SQLException se) {  
    System.out.println("Error: " + se);  
}
```

Execute
Query and
store results
in ResultSet

Get value in
first column

JDBC: Iterate result set

```
try {  
    ResultSet resultSet = statement.executeQuery("select  
        name, room from Professor where rank = 'C4'");  
    System.out.println("C4-Professors:");  
    while (resultSet.next()) {  
        System.out.println(resultSet.getString("Name") +  
            " " + resultSet.getInt("Room"));  
    }  
    resultSet.close();  
} catch (SQLException se) {  
    System.out.println("Error: " + se);  
}
```

Execute
Query and
store results
in ResultSet

Iterate over
resultSet



JDBC: Close connection

```
try {  
    statement.close();  
    connection.close();  
} catch (SQLException e) {  
    System.out.println("Error when closing DB connection: "  
        + e);  
}
```

Close statement
and connection

Exercise 3: JDBC

Get started:

- Download the project from moodle
- Run the DatabaseCreator
- You can look at the database either by
 - Double-clicking on the h2-1.4.179.jar (Only if you associated .jar files with Java)
 - Run the jar from the terminal using `java -jar h2*.jar`

Task:

- Formulate an SQL query for each of the tasks below
- Execute the query using java code
- Log the result

Queries:

- Count the number of C4 professors.
- Select all students with less than 8 semesters

JDBC: Prepared Statements

- Creates a Statement, which is easy to reuse with different parameters
e.g. Finding several students by name
- Faster execution of batches
- Checks parameters, hence it helps avert SQL injections

JDBC: Prepared Statements

```
String statementString = "update contact set  
firstname = ? where contactID = ?";
```

? is placeholder for
concrete values

String containing
query

```
PreparedStatement updateStatement =  
con.prepareStatement(statementString );
```

Create
PreparedStatement

```
updateStatement.setString(1, "Dieter");  
updateStatement.setInt(2, 5);
```


Index of
placeholder

Value

Replace placeholders
with values

JDBC: Prepared Statements

User Input is always a potential security hazard!



Drop table removes a table

ORM (Object-Relational-Mapping)

- Directly maps objects to a relational database → Virtual object database
- Advantage:
 - Save code
 - Work on objects (higher abstraction)
- Disadvantage:
 - Potential performance hit due to high level of abstraction
 - Might create poorly designed databases

ORM Mapping for one class: Entity


Student
- name: String - matrNumber: int - semester: int
+ getGradeAverage() : int + sumSemesterHours() : short

Mapping can easily done using **annotations**

```
@Entity
public class Student {
    @Id
    @GeneratedValue(strategy=GenerationType.AUTO)
    private int matrNumber;
    ...
}
```

ORM Mapping of relations between classes

- Also done using annotations: @OneToOne, @ManyToOne, @ManyToMany



```
@Entity
public class Student {
    @Id
    @GeneratedValue(strategy = GenerationType.AUTO)
    private int matrNumber;

    @ManyToMany(fetch=FetchType.LAZY)
    private List<Course> courses;
    ...
}
```

ORM: Entity Manager

Create new Student

```
EntityManager entityManager =  
  
entityManagerFactory.createEntityManager();  
entityManager.getTransaction().begin();  
entityManager.persist( new Student(12354,  
    "Donald", 12);  
entityManager.getTransaction().commit();  
entityManager.close();
```

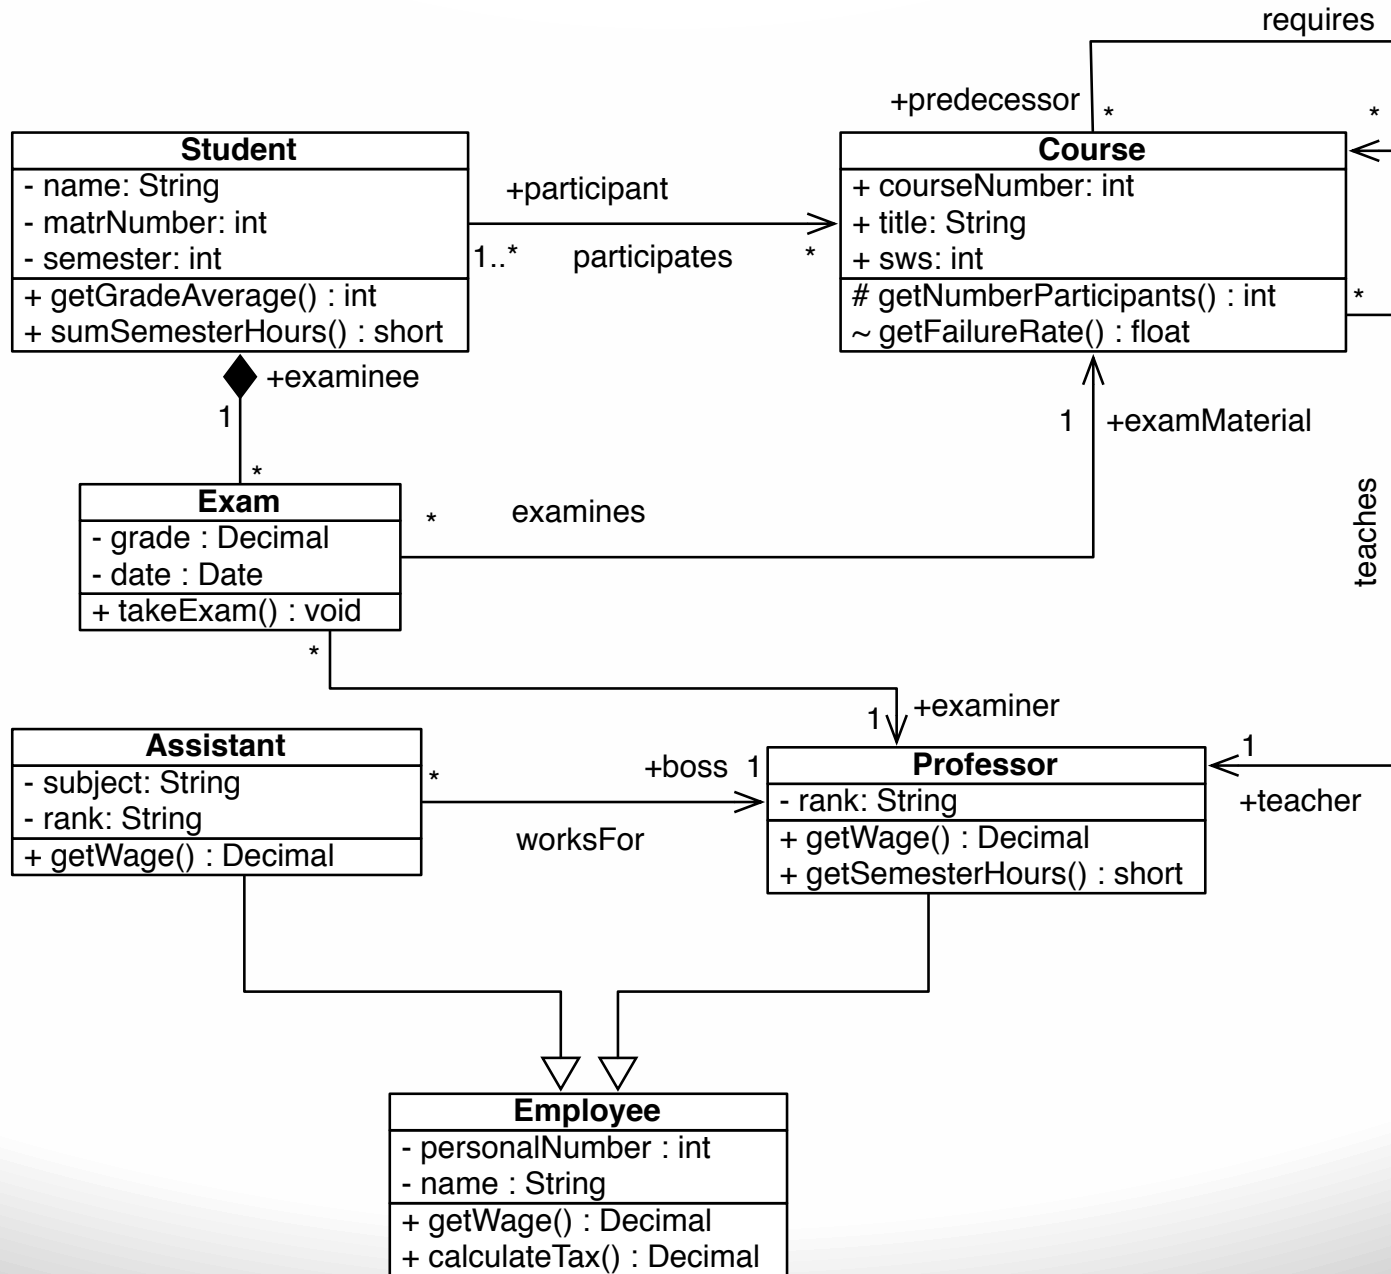

ORM: Entity Manager

Find a student

```
entityManager =
    entityManagerFactory.createEntityManager();
entityManager.getTransaction().begin();
List<Student> result = entityManager.createQuery
( "from
    Student", Student.class ).getResultList();
for ( Student student : result ) {
    System.out.println( "Student (" + student.get
        Name() + ")");
}
entityManager.getTransaction().commit();
entityManager.close();
```

Exercise 4:

- Write the professor class
- Add annotations to support ORM
(`@Entity`, `@Id`, `@OneToOne`, `@OneToMany`, `@ManyToMany`)
- Assume that all other classes are provided
- Professor should have `personalNumber`, `name`, and `rank`



NoSQL (Not only SQL)

- Motivation (Why NoSQL)
 - Scalability
- Solution: Horizontal Scaling
- Major use case: Interactive web applications
 - Examples:
 - Social Networks (Facebook, Twitter, Digg)
 - Auction Platforms (eBay)
- Problems
 - Only weak support for consistency
 - Lack of standardized interfaces
 - Only low-level query languages

Overview

- Databases have several goals:
 - Avoid redundancy and inconsistency
 - Rich access to data
 - Synchronize concurrent data access
 - Avoid loss of data
 - Enforcing integrity rules
 - Security and Privacy
- There exist several types of databases, most important
 - Relational database
 - Objects database
 - NoSQL database
- Relational databases
 - Consist of tables
 - Tables are defined using a schema, which includes data types and integrity information
 - **SQL** is a language that allows to create, alter and query a relational database
- ORM
 - Maps objects into a relational database automatically
 - Usually implemented using annotations or subclassing
- NoSQL
 - Attempt to deal with performance issues in huge real-time web applications
 - Often sacrifice consistency to achieve higher performance

Bibliography

